

Interrogation de données structurées en Spark

Master DAC – Bases de Données Large Echelle

Mohamed-Amine Baazizi

mohamed-amine.baazizi@lip6.fr

2021-2022

High-level data pipelines

- The RDD algebra is not well-suited
 - Unfriendly programming, unreadable code
 - Misses the opportunity of logical optimization
- Most data pipelines contain
 - Analytical queries
 - Well-defined transformations (ex. aggregation, *feature extraction*)
- Most data is structured and have an implicit schema
- Adopt a declarative approach
 - SQL queries
 - Dataframe expressions: inspired from Data Science libraries
- The RDD algebra serves as a physical layer

Spark Dataframe

- Abstraction on top of RDDs
 - RDD[Row], Row = sequence of values with a fixed schema
- Several operators
 - Relational operator
 - User Defined and built-in functions
 - Maps each element to another element
 - E.g. in IMBD, extract year from title based on a regexp
- Distinction between transformations and actions
 - Some useful actions
 - `printSchema()`, `show()`, `read()`, `write()`
- Support for different formats
 - Tabular (CSV, TSV), nested (JSON, Parquet), Text
 - Specific formats (libsvm for ML)

Data model

- Base types
 - boolean, numeric family (integer, decimal, ...), String, null, timestamp
- Complex types
 - Arrays : (type of element, containsNull) → *homogenous*
 - Structure :
 - [StructField, ..., StructField]
 - StructField(name, type, nullable) → name is a string and is unique
 - Maps:
 - (keyType, valueType, valueContainsNull) → key of any type and is unique
- Support for user-defined types (UDT)
 - Application-specific data
 - Object-oriented databases style
 - Data and methods Dataset
 - Dataset with Scala

containsNull, nullable and valueContainsNull indicate the presence of null, but are always set to true

Tabular data

```
movies = spark.read.format("csv"). ...
```

```
movies.show(truncate=False)
```

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance

```
movies.printSchema()
```

```
root
|-- movieId: long (nullable = true)
|-- title: string (nullable = true)
|-- genres: string (nullable = true)
```

Useful Dataframe operators

- Relational algebra
 - Unary : select, where
 - Binary : join, intersect, subtract, union
- Schema
 - drop: removes a column
 - withColumn: adds a column
 - withColumnRenamed: renames a column
- Grouping
 - groupBy(col*)
 - groups on a set of columns
 - Produces a *GroupedDataset*
- Sorting
- Aggregation

GroupedDataset operators

- Partition-wise single aggregation
 - `min(col)`, `max(col)`, `sum(col)`, `count()`
 - return type = Dataframe with one column
- Partition-wise multiple aggregations
 - `agg(min(col), max(col), ...)`
 - return type = Dataframe with many columns
- Pivoting
 - `Pivot(col) -> Dataframe(rcol_1, ..., rcol_n)`
 - For `i` in `[1, #values(col)]` : `name(rcol_i) = value(col_i)`

Built-in functions

- Array/maps manipulation
 - Arrays: containment, distinct, intersect, except, max/min, search, size, sort,...
 - Maps: concatenation, filtering, size,...
- General-purpose
 - Math: descriptive stats, trigonometry, skewness
 - Summary: *countDistinct*, *sumDistinct*, ...
 - Elements to Collections and vice-versa
 - Nesting: *collect_list*, *collect_set*
 - Flattening: *explode*, *explode_outer*
 - Temporal:
 - current date, date arithmetic, day/month/year extraction, conversion
 - Data transformation
 - Json/csv parsing
 - Timestamp extraction/encoding
 - Strings: length, distance, trim, regexp extraction , splitting, case conversion, ...
- Fully-documented online

Illustration:

Building and flattening lists

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance

```

root
|-- movieId: long
|-- title: string
|-- genres: string

```

```
from pyspark.sql.functions import col, split, explode
```

```
genres_list = movies.select(split(movies.genres, '|').alias('genre_list'))
```

Size → 7

```
genres = genres_list.select(explode(col('genre_list')).alias('genre')).distinct()
genres.count()
```

10

genre
Adventure
Comedy
Romance
Children

Illustration:

Pivoting rows to columns

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance

```

root
|-- movieId: long
|-- title: string
|-- genres: string

```

```

movies_genre = movies.select('movieId',explode(split(movies.genres,'\\|')).alias('genre'))\
.groupBy('movieId').pivot('genre').count().orderBy('movieId')

```

movieId	Action	Adventure	Animation	Children	Comedy	Crime	Drama	Fantasy	Romance	Thriller
1	null	1	1	1	1	null	null	1	null	null
2	null	1	null	1	null	null	null	1	null	null
3	null	null	null	null	1	null	null	null	1	null

...

Illustration: applying built-in functions

userId	movieId	rating	timestamp
1	1	2.5	1260759144
1	2	3.0	1260759179
2	1	3.0	1260759182
4	3	2.0	1260759185

```
userId: long (nullable = true)
movieId: long (nullable = true)
rating: double (nullable = true)
timestamp: long (nullable = true)
```

```
from pyspark.sql.functions import from_unixtime, year, month, dayofmonth
```

```
ratings_date = ratings.select(from_unixtime('timestamp').alias('datetime'))\
.select('datetime', dayofweek('datetime').alias('day'))
```

datetime	day
2009-12-14 02:52:24	2
2009-12-14 02:52:59	2
2009-12-14 02:53:02	2
2009-12-14 02:53:05	2

User-defined Functions

- Express non-relational operations
 - Ex. map rating to categories, compute user similarity
 - Supported in most database systems
 - Executed out of the SQL context
 - No automatic optimisation but
 - Possibility of adding a new optimisation rules (for experts)
 - If not carefully designed, performance degradation
 - Only use when no available built-in function can do the work
 - Two possibilities in Spark
 - Standard: invoke on each row → costly
 - Vectorized: invoke on a vector (batch of rows) → more efficient
- Benefit from memory columnar storage (Arrow format)

Row-at-time vs vectorized UDFs

col1	col2	...
v00	v01	
v10	v11	
v20	v21	

std(v00)
std(v10)
std(v20)

```
@udf()  
def std(): ...  
select(std(col1))
```

col1	col2	...
v00	v01	
v10	v11	
v20	v21	

Panda([v00,v10,v20])

```
@pandas_udf()  
def panda(): ...  
select(panda(col1))
```

panda() must preserve the input size
and be independent of the data partitioning

Illustration: applying a standard UDF

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance

root

- movieId: long
- title: string
- genres: string

```
from pyspark.sql.functions import udf
```

```
@udf('string')
def nb_genres(l):
    return str(len(l))+' genres(s)'
```

```
select('movieId', 'title', nb_genres('genre_list').alias('nb_genres'))
```

movieId	title	nb_genres
1	Toy Story (1995)	5 genres(s)
2	Jumanji (1995)	3 genres(s)
3	Grumpier Old Men (1995)	2 genres(s)

Illustration: applying a vectorized UDF

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance

root
|-- movieId: long
|-- title: string
|-- genres: string

```
import pandas as pd
from pyspark.sql.functions import col, pandas_udf
from pyspark.sql.types import ArrayType
```

```
def nb_genres(s:pd.Series) -> pd.Series:
    return s.str.split('\|')
```

```
panda_nb_genres = pandas_udf(nb_genres, returnType=ArrayType(StringType()))
```

```
movies.select('movieId', 'title', panda_nb_genres('genres').alias('nb_genres'))
```

```
root
|-- movieId: long (nullable = true)
|-- title: string (nullable = true)
|-- nb_genres: array (nullable = true)
|   |-- element: string (containsNull = true)
```

Data reading and writing

- Reading / Writing
 - Several sources (local files, HDFS), batch or streaming
 - Many database connectors: JDBC or system-specific (eg. Mongo, Couchbase)
 - Guided by a schema
 - Schema is either provided or inferred
- Reading/writing options
 - Fully documented (see API reference)
 - Format, encoding, separators, compression
- Side effect free by default
 - Modify a copy of the data
- Potential solution: Delta Lake
 - ACID support: insert, Update, delete operations
 - Data versioning and schema evolution
 - Metadata management

Schema management

- Option 1: passed as argument
 - Plain text
 - 'attr_name type, ..., attr_nametype'
 - *Dataframe* object
 - StructType(), ArrayType(), LongType(), ...
 - Basic schema enforcement
 - Type compatibility, Nullity
 - Failure raises an exception

```
f1 = StructField('f1', IntegerType(), True)
f2 = StructField('f2', IntegerType(), True)
df_sch = StructType([f1,f2])
```

```
df = spark.createDataFrame(data, df_sch)
```

data: RDD with [[1,1],[None,2], [3,None]]

```
root
|-- f1: integer (nullable = true)
|-- f2: integer (nullable = false)
```



Schema management

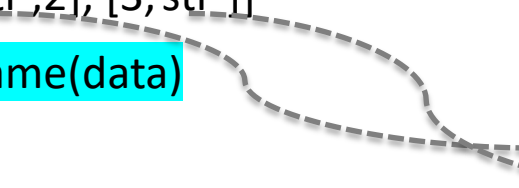
- Option 2: inferred from the data
 - Sample of the data, usually 10%
 - Useful when not shipped with the data, or when data is nested (JSON)
 - Type promotion rules
 - Triggered to solve conflicting types
 - Lossy transformation

data: RDD with [[1,1],['str',2], [3,'str']]

```
df = spark.createDataFrame(data)
```

root

```
|-- _1: long (nullable = true)
|-- _2: long (nullable = true)
```



_1	_2
1	1
null	2
3	null

Beyond Spark schema management

- Schema enforcement strategy
 - Only valid data or mix between non-valid and valid?
 - Discard tuples, substitute with default?
- Support integrity constraints (IC)
 - Primary key - foreign key
 - Domain constraints
 - Assertions
- Data quality issues
 - Missing values
- Potential solution: Delta Live Tables
 - Automatic support for schema evolution
 - IC specification and validation
 - Metadata collection at scale
 - Hassle-free ETL

Managing nested data

- Schema inference, data loading
 - Automatic, on a sample of the data
 - Data is shredded using the schema
 - Shredding rules
 - JSON Object → Dataframe Struct
 - JSON Arrays → Dataframe Array
 - JSON Basic types → corresponding Dataframe basic types
 - Missing fields indicated with null
 - Conflict resolution using type promotion rules

JSON and irregularity

Input: 2 documents

```
{
  "person": {
    "firstname": "Melena",
    "lastname": "RZYIK",
    "role": "reported",
    "rank": 1,
    "organization": "abc"
  }
}
```

```
{
  "person": {
    "firstname": "other",
    "lastname": "ABCD",
    "rank": 1,
    "organization": "OO"
  }
}
```

```
spark.read.json()
```

root

```
-- person: struct (nullable = true)
|   |-- firstname: string (nullable = true)
|   |-- lastname: string (nullable = true)
|   |-- organization: string (nullable = true)
|   |-- rank: long (nullable = true)
|   |-- role: string (nullable = true)
```

schema

firstname	lastname	organization	rank	role
Melena	RZYIK	abc	1	reported
other	ABCD	OO	1	

Missing role

data

JSON and irregularity

Input: 3 documents

```
{  
  "first" : "al",  
  "coord" : [],  
  "last" : "jr"  
}
```

```
{  
  "first" : "al",  
  "coord" : null,  
  "last" : "jr"  
}
```

```
{  
  "email" : "abc@ef",  
  "first" : "li",  
  "coord" : {  
    "lat" : 45,  
    "long" : 12  
  },  
  "last" : null  
}
```

```
spark.read.json()
```

root

- coord: string (nullable = true)
- email: string (nullable = true)
- first: string (nullable = true)
- last: string (nullable = true)

coord	email	first	last
[]	null	al	jr
null	null	al	jr
{"long":12,"lat":45}	abc@ef	li	null

schema

Object stored as a string!!

data

Querying nested data

- Navigating the hierarchy of objects
 - Dot-notation
 - `select (' level1.level2.)`
 - Traversing arrays is not allowed using the dot-notation!!
- Accessing the content of arrays
 - Flatten then navigate
 - `explode()`, `explode_outer()`
 - Use existing (built-in) array functions
 - E.g. `array_exist()`, `array_contains()`
 - Performance issues:
 - Flattening allows for standard logical optimization to be applied
 - Many research effort on how to rewrite queries to take advantage of logical optimization

Navigating structures

```
{
  "person" : {
    "firstname" : "Melena",
    "lastname" : "RYZIK",
    "role" : "reported",
    "rank" : 1,
    "organization" : ""
  }
}
```

```
{
  "person" : {
    "firstname" : "other",
    "lastname" : "ABCD",
    "rank" : 1,
    "organization" : "OO"
  }
}
```

```
root
|-- person: struct (nullable = true)
|   |-- firstname: string (nullable = true)
|   |-- lastname: string (nullable = true)
|   |-- organization: string (nullable = true)
|   |-- rank: long (nullable = true)
|   |-- role: string (nullable = true)
```

select("person.*")

firstname	lastname	organization	rank	role
Melena	RYZIK		1	reported
other	ABCD	OO	1	null

Accessing Arrays

```
{
  "person" : [
    {
      "firstname" : "Melena",
      "lastname" : "RYZIK",
      "role" : "reported",
      "rank" : 1,
      "organization" : ""
    },
    {
      "firstname" : "derba",
      "lastname" : "OKYZ",
      "role" : "reported",
      "rank" : 1,
      "organization" : ""
    }
  ]
}
```

```
{
  "person" : [
    {
      "firstname" : "other",
      "lastname" : "ABCD",
      "rank" : 1,
      "organization" : "OO"
    }
  ]
}
```

`select(col("person"),explode(col("person")))`

root

```
|-- person: struct (nullable = true)
|   |-- firstname: string (nullable = true)
|   |-- lastname: string (nullable = true)
|   |-- organization: string (nullable = true)
|   |-- rank: long (nullable = true)
|   |-- role: string (nullable = true)
```

person	col
[[Melena, RYZIK, , 1, reported], [derba, OKYZ, , 1, reported]]	[Melena, RYZIK, , 1, reported]
[[Melena, RYZIK, , 1, reported], [derba, OKYZ, , 1, reported]]	[derba, OKYZ, , 1, reported]
[[other, ABCD, OO, 1,]]	[other, ABCD, OO, 1,]

Closing remarks

- Dataframes
 - a full-fledged API for building complex ETL and analytics
 - many room for physical optimization
 - Disk and memory columnar storage (following lectures), indexing
- Interesting topics
 - The Delta Lake approach
 - Code versioning
- 2 lab sessions
 - Querying Tabular data: gain familiarity with the Dataframe algebra, use functions
 - Querying Nested data: deal with irregularity and incompleteness, real-world data
- *Next 30': Zennea presentation on how to manage metadata at an enterprise level*